

# STAT 6340 Mini Project 4

Matthew Lynn  
March 27, 2019

## Section 1

Section 2 is coded in Section 1

### Question 1(a)

We make a scatterplot of gpa against act.  
We notice there is a positive correlation but it is not strong (0.27)

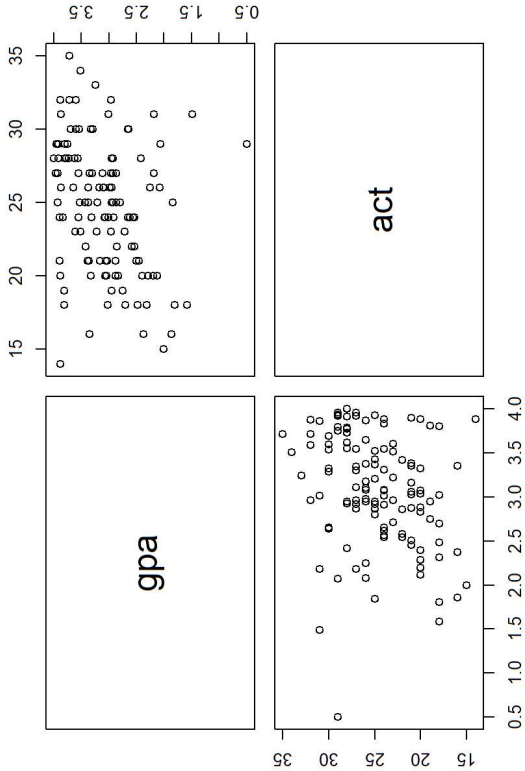
```
gpa = read.csv("gpa.csv", header = T)
str(gpa)
```

```
## 'data.frame': 120 obs. of 2 variables:
## $ gpa: num 3.9 3.88 3.78 2.54 3.03 ...
## $ act: int 21 14 28 22 21 31 32 27 29 26 ...
```

```
summary(gpa)
```

	gpa	act
## Min.	:0.500	Min. :14.00
## 1st Qu.	:2.689	1st Qu.:21.00
## Median	:3.078	Median :25.00
## Mean	:3.074	Mean :24.73
## 3rd Qu.	:3.593	3rd Qu.:28.00
## Max.	:4.000	Max. :35.00

```
pairs(subset(gpa))
```



```
round(cor(subset(gpa)), 2)
```

```
##      gpa  act
## gpa 1.00 0.27
## act 0.27 1.00
```

```
# There is a positive correlation between gpa and act
# the relationship is not very strong (0.27)
```

### 1(b)

$\rho$  is our population correlation between gpa and act.  
Here, we gather the bootstrap estimates of bias and standard error of the point estimate and show the 95% confidence interval from bootstrap

```
library(boot)
corr(gpa)
```

```
## [1] 0.2694818
```

# 1(c)

Using gpa~act, we fit a SLR. We show the least square estimates, standard errors and a 95% CI. We include verifications of our model assumptions

```
fit = lm(gpa~act, data = gpa)
# Least square estimate of coefficients, SE, and 95% CI
summary(fit)
```

```
## Call:
## lm(formula = gpa ~ act, data = gpa)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.74004 -0.33827  0.04062  0.44064  1.22737
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.11405    0.32089   6.588 1.3e-09 ***
## act         0.03883    0.01277   3.040 0.00292 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6231 on 118 degrees of freedom
## Multiple R-squared:  0.07262, Adjusted R-squared:  0.06476
## F-statistic: 9.24 on 1 and 118 DF, p-value: 0.002917
```

```
confint(fit)
```

```
##              2.5 %      97.5 %
## (Intercept) 1.47859015 2.74950842
## act         0.01353307 0.06412118
```

```
par(mfrow = c(2,2))
plot(fit)
```

```
# either function below will work
corr.fn <- function(data, i=c(1:length(data))) {
  result <- data[i,]
  return(corr(result$gpa, result$act))
}
corr.fn <- function(data, i=c(1:length(data))) {
  result = corr(data[i,])
  return(result)
}
set.seed(1)
corr.boot <- boot(gpa, corr.fn, R = 1000)
# Point estimate, bootstrap bias, and bootstrap se
corr.boot
```

```
## ORDINARY NONPARAMETRIC BOOTSTRAP
##
## Call:
## boot(data = gpa, statistic = corr.fn, R = 1000)
##
## Bootstrap Statistics :
##      original    bias      std. error
## t1* 0.2694818 0.007800885  0.1072831
```

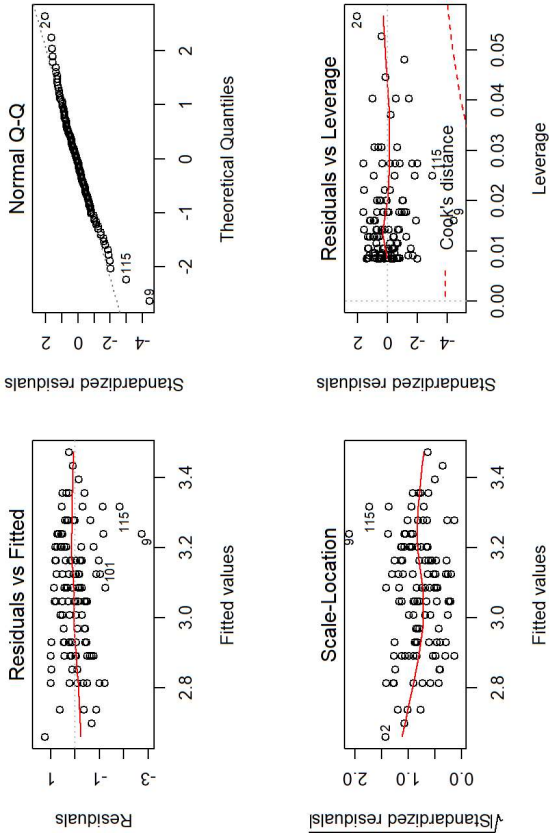
```
# 95% bootstrap conf int
boot.ci(corr.boot, type = "perc")
```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
## CALL :
## boot.ci(boot.out = corr.boot, type = "perc")
##
## Intervals :
## Level      Percentile
## 95%      ( 0.0728,  0.4917 )
## Calculations and Intervals on Original Scale
```

```
# verified below
sort(corr.boot$t)[c(25, 975)]
```

```
## [1] 0.07274536 0.49136740
```

# our original p is contained in the bootstrap interval and the bias is only # 0.004196114 which implies good bootstrap estimates.



```
par(mfrow = c(1,1))
# it appears the residual plot could do better via a transformation but not bad
# the qq plot is not bad but does show some outliers.
```

## 1(d)

We use a nonparametric bootstrap to estimate the previous question's parameters. The bootstrap estimates are higher than was in the linear model. Also, the 95% CI is wider than the linear model.

```
library(boot)
# function to output coefficients from linear model
fit.fn <- function(data, index) {
  result <- coef(lm(gpa ~ act, data = gpa, subset = index))
  return(result)
}
set.seed(1)
fit.boot = boot(gpa, fit.fn, R = 1000)
fit.boot
```

```
## ORDINARY NONPARAMETRIC BOOTSTRAP
##
## Call:
## boot(data = gpa, statistic = fit.fn, R = 1000)
##
## Bootstrap Statistics :
##      original      bias      std. error
## t1* 2.11404929 -0.0108361368  0.35962727
## t2* 0.03882713  0.0004350197  0.01455027
```

```
boot.ci(fit.boot, type = "perc", index = 1) # intercept
```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = fit.boot, type = "perc", index = 1)
##
## Intervals :
## Level      Percentile
## 95%      ( 1.387,  2.812 )
## Calculations and Intervals on Original Scale
```

```
boot.ci(fit.boot, type = "perc", index = 2) # act
```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 1000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = fit.boot, type = "perc", index = 2)
##
## Intervals :
## Level      Percentile
## 95%      ( 0.0110,  0.0674 )
## Calculations and Intervals on Original Scale
```

## Question 2(a)

We want to examine StoreID, STORE and Store7 variables from the OJ data. After a short EDA we determine that StoreID seems to contain all the info in STORE and Store7. Thus, we can drop STORE and Store7. We then use Purchase as our response, take StoreID as categorical, and split our data into a train set and test set for training and predicting.

```
library(ISLR)
names(OJ)

## [1] "Purchase"      "WeekofPurchase" "StoreID"      "PriceCH"
## [5] "PriceMW"       "DiscH"          "DiscMW"       "SpecialCH"
## [9] "SpecialMW"    "LoyalCH"       "SalePriceMW"  "SalePriceCH"
## [13] "PriceDiff"    "Store7"        "PctDiscMW"   "PctDiscCH"
## [17] "ListPriceDiff" "STORE"


# use only these 4 variables
OJ = subset(OJ, select = c(Purchase, StoreID, Store7, STORE))
str(OJ)

## 'data.frame':  1070 obs. of  4 variables:
## $ Purchase: Factor w/ 2 levels "CH","MW": 1 1 1 2 1 1 1 1 1 1 ...
## $ StoreID : num  1 1 1 7 7 7 7 7 ...
## $ Store7  : Factor w/ 2 levels "No","Yes": 1 1 1 2 2 2 2 2 ...
## $ STORE   : num  1 1 1 0 0 0 0 0 ...

# STORE is a subset of StoreID and Store7 is transitive through StoreID
pairs(OJ)
```

```
fit = glm(Purchase~., family = binomial, data = OJ)
summary(fit)

## Call:
## glm(formula = Purchase ~ ., family = binomial, data = OJ)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.2736  -0.9733  -0.7236   1.1852   1.7136
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.46471    0.19347   2.402  0.0163 *
## StoreID     -0.24147    0.07339  -3.290  0.0010 **
## StoreYes     0.01916    0.36632   0.052  0.9583
## STORE       NA          NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## (Dispersion parameter for binomial family taken to be 1)
##      Null deviance: 1430.9  on 1069  degrees of freedom
## Residual deviance: 1360.4  on 1067  degrees of freedom
## AIC: 1366.4
##
## Number of Fisher Scoring iterations: 4
```



```
rm(OJ)
# It appears that only StoreID and factor Level "No" (intercept) from
# Store7 have any significance when all predictors are used.
# we drop Store7 and STORE
newOJ = subset(OJ, select = c(-Store7, -STORE))
# I do not see any reason to force MM = 1 and CH = 0
# When I experimented, I was able to produce the same results
newOJ$StoreID = as.factor(newOJ$StoreID) # take StoreID as factor
# str(newOJ) # uncomment to verify is factor
# split data 50/50, train/test
set.seed(1)
n = nrow(newOJ)
sampler = sample(1:n, n/2) # n/2 is the 50/50 splitter
train = newOJ[sampler, ]
test = newOJ[-sampler, ]
```

## 2(b)

Here we train a logistic regression model and use the confusion matrix to obtain sensitivity, specificity, overall misclassification, then plot the ROC curve and estimate using a 10 fold cross validation

The caret package does all of this nicely and compactly, so we proceed

through the next few questions using similar coding.

The accuracy from our model (train data) has accuracy (1-test error) very close to our confusion matrix accuracy (test data, predicted)

The ROC curve has a nice left corner shape, notice that the graph continues to the left till -0.5, needs to be fixed to origin preferably (sorry!) pROC shapes it this way with `asp = 1` to keep the shape square


```
# 10 fold cv
# Library(crossval) # could not make sense of its inner workings
# Even their examples needed tailoring, and even after that
# it was not very intuitive. I instead proceed with Caret package.
# When using K-fold CV, we use the original data set unsplit.
# However, I split the data beforehand so I could make a predict set
# with the test data. But in general this package's function should
# do its own K splits as it trains/tests on the data.
# A rather robust way of estimating accuracy (or test error rate)
library(caret)
set.seed(1)
tc <- trainControl(method = "cv", number = 10, verboseIter = F)
model <- train(Purchase~., train,
               method="glm", family="binomial", trControl = tc)
model
```

```
## Generalized Linear Model
##
## 535 samples
## 15 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 481, 482, 481, 481, 482, 482, ...
## Resampling results:
##
## Accuracy Kappa
## 0.8204403 0.6202817
```

```
# 10-fold CV estimated test error rate = 1 - overall misclassification
# also, this is 1 - accuracy = test error rate (estimated)
glm.err = 1 - model$results[,2]
prob = predict(model, test, type = "prob")
# pred = predict(model, test) # or
pred <- ifelse(prob$MM >= 0.5, "MM", "CH")
con.mat = table(test[, "Purchase"], pred)
confusionMatrix(con.mat)
```

```
## Confusion Matrix and Statistics
##
##      pred
##      CH  MM
## CH 287  43
## MM  53 152
##
##      Accuracy : 0.8206
##      95% CI : (0.7854, 0.8522)
##      No Information Rate : 0.6355
##      P-Value [Acc > NIR] : <2e-16
##
##      McNemar's Test P-Value : 0.3583
##
##      Sensitivity : 0.8441
##      Specificity : 0.7795
##      Pos Pred Value : 0.8697
##      Neg Pred Value : 0.7415
##      Prevalence : 0.6355
##      Detection Rate : 0.5364
##      Detection Prevalence : 0.6168
##      Balanced Accuracy : 0.8118
##
##      'Positive' Class : CH
##
```

```
library(pROC)
roc <- roc(test[, "Purchase"], prob$MM, levels = c("MM", "CH"))
plot(roc, legacy.axes = T)
```



## 2(c)

We employ the same methods but using LDA instead


```
library(caret)
set.seed(1)
model <- train(Purchase~., train,
               method="lda", trControl = tc)
model
```

```
## Linear Discriminant Analysis
##
## 535 samples
## 15 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 481, 482, 481, 481, 482, 482, ...
## Resampling results:
##
## Accuracy      Kappa
## 0.8205101 0.6215449
```

```
lda.err = 1 - model$results[,2]
prob = predict(model, test, type = "prob")
pred <- ifelse(prob$MM >= 0.5, "MM", "CH")
confusionMatrix(table(test$Purchase, pred))
```

```
## Confusion Matrix and Statistics
##
##      pred
##      CH  MM
## CH 285  45
## MM  52 153
##
##          Accuracy : 0.8187
##          95% CI : (0.7834, 0.8504)
##        No Information Rate : 0.6299
##        P-Value [Acc > NIR] : <2e-16
##
##        Mcnemar's Test P-Value : 0.6139
##
##          Sensitivity : 0.8457
##          Specificity : 0.7727
##        Pos Pred Value : 0.8636
##        Neg Pred Value : 0.7463
##        Prevalence : 0.6299
##        Detection Rate : 0.5327
##        Detection Prevalence : 0.6168
##        Balanced Accuracy : 0.8092
##
##        'Positive' Class : CH
##
```

```
library(pROC)
roc <- roc(test$Purchase, prob$MM, levels = c("MM", "CH"))
plot(roc, legacy.axes = T)
```



## 2(d)

We employ the same methods but using QDA instead QDA fails if there is rank deficiency in other words we have columns that fully or partially (near fully) explain other columns. This means our data has linear dependency or multicollinearity. If we can't invert a matrix then we cannot use qda effectively or at all. Below is what would be the qda procedure if there were no multicollinearity. However, it will result in rank deficiency and ultimately fail.

```
# Library(caret)
# set.seed(1)
# model <- train(Purchase~, train,
#               method="qda", trControl = tc)
# model
qda.err = 1 - model$results[,2]
# prob = predict(model, test, type = "prob")
# pred <- ifelse(prob$MM >= 0.5, "MM", "CH")
# confusionMatrix(table(test$Purchase, pred))
# Library(pROC)
# roc <- roc(test$Purchase, prob$MM, levels = c("MM", "CH"))
# plot(roc, legacy.axes = T)
```

## 2(e)

We employ the same methods but using KNN, we also use caret to find the optimal KNN k value for lowest test error rate


```
library(caret)
set.seed(1)
model <- train(Purchase~, train,
               method="knn", trControl = tc)
model # optimal k is 5, caret automatically uses this k

## k-Nearest Neighbors
##
## 535 samples
## 15 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 481, 482, 481, 482, 482, ...
## Resampling results across tuning parameters:
##
## k Accuracy Kappa
## 5 0.6895178 0.3432383
## 7 0.6743885 0.3045327
## 9 0.6633124 0.2722348
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 5.

knn.err = 1 - model$results[,2]
prob = predict(model, test, type = "prob")
pred <- ifelse(prob$MM >= 0.5, "MM", "CH")
confusionMatrix(table(test$Purchase, pred))
```

```
## Confusion Matrix and Statistics
##
##      pred
##      CH  MM
## CH 262  68
## MM  89 116
##
##      Accuracy : 0.7065
##      95% CI : (0.666, 0.7448)
##      No Information Rate : 0.6561
##      P-Value [Acc > NIR] : 0.007396
##
##      Kappa : 0.3669
##      Mcnemar's Test P-Value : 0.110450
##
##      Sensitivity : 0.7464
##      Specificity : 0.6304
##      Pos Pred Value : 0.7939
##      Neg Pred Value : 0.5659
##      Prevalence : 0.6561
##      Detection Rate : 0.4897
##      Detection Prevalence : 0.6168
##      Balanced Accuracy : 0.6884
##
##      'Positive' Class : CH
##
```

```
library(ROC)
roc <- roc(test$purchase, prob$MM, levels = c("MM", "CH"))
plot(roc, legacy.axes = T)
```



2(f)

Now lets combine all the test error rates together and see which is the lowest  
It seems that lda is our best model for 10 fold CV  
KNN had nearly double of the other models

```
rbind(glm.err, lda.err, qda.err, knn.er = min(knn.err))
```

```
##
##      [,1]
## glm.err "0.179559748427673"
## lda.err "0.17948986725716"
## qda.err "No good"
## knn.er  "0.310482180293501"
```

Question 3(a)

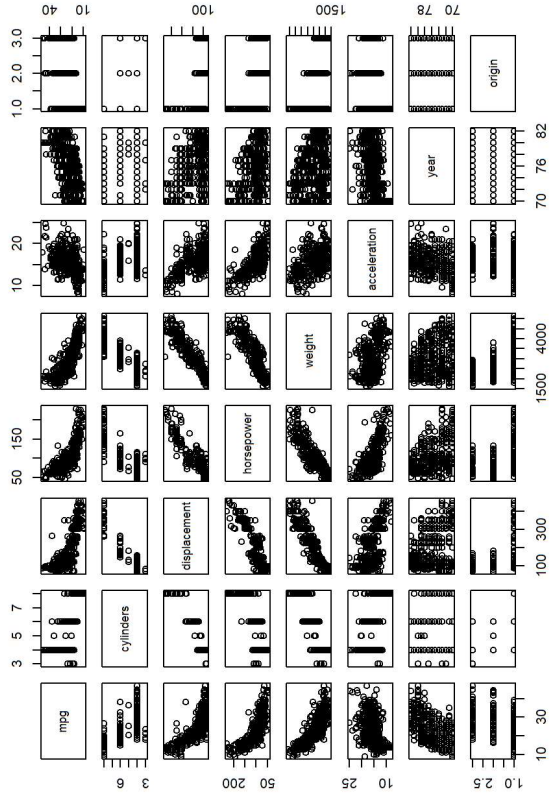
Let us dive into the Auto data from ISLR and use MPG as our response



```
library(ISLR)
library(leaps)
set.seed(1)
auto = subset(Auto, select = -name)
# eda
str(auto)
```

```
## 'data.frame': 392 obs. of 8 variables:
## $ mpg : num 18 15 18 16 17 15 14 14 14 15 ...
## $ cylinders : num 8 8 8 8 8 8 8 ...
## $ displacement: num 307 350 318 304 302 429 454 440 455 390 ...
## $ horsepower : num 130 165 150 150 140 198 220 215 225 190 ...
## $ weight : num 3504 3693 3436 3433 3449 ...
## $ acceleration: num 12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
## $ year : num 70 70 70 70 70 70 70 70 70 ...
## $ origin : num 1 1 1 1 1 1 1 1 1 ...
```


```
# summary(auto) # nice way to get some quick stats, but messy so I commented
# Just Looking at the scatterplots we see that mpg has many trends with others
pairs(subset(auto))
```



```
cor(auto)
```

```
## mpg cylinders displacement horsepower weight
## mpg 1.000000 -0.7776175 -0.8051269 -0.7784268 -0.8322442
## cylinders -0.7776175 1.0000000 0.9508233 0.8429834 0.8975273
## displacement -0.8051269 0.9508233 1.0000000 0.8972570 0.9329944
## horsepower -0.7784268 0.8429834 0.8972570 1.0000000 0.8645377
## weight -0.8322442 0.8975273 0.9329944 0.8645377 1.0000000
## acceleration 0.4233285 -0.5046834 -0.5438005 -0.6891955 -0.4168392
## year 0.5805410 -0.3456474 -0.3698552 -0.4163615 -0.3091199
## origin 0.5652088 -0.5689316 -0.6145351 -0.4551715 -0.5850054
## mpg cylinders displacement horsepower weight
## mpg 0.4233285 0.5805410 0.5652088
## cylinders -0.5046834 -0.3456474 -0.5689316
## displacement -0.5438005 -0.3698552 -0.6145351
## horsepower -0.6891955 -0.4163615 -0.4551715
## weight -0.4168392 -0.3091199 -0.5850054
## acceleration 1.0000000 0.2903161 0.2127458
## year 0.2903161 1.0000000 0.1815277
## origin 0.2127458 0.1815277 1.0000000
```

```
# seems to be right skewed
hist(auto$mpg)
```



```
# here we look at what would be worthy predictor variables. However,
# we are using ALL of them and whittling down with variable selection methods.
```

### 3(b)

For this part we want to use MLR using least squares  
To make this part short we simply fit all predictors as our full model  
Then run a summary and refit using only the significant ones  
Lastly, we compare the models using partial F test and see that  
the reduced model is better.

```
# fit a MLR
fit = lm(mpg~., data = auto)
summary(fit)

## Call:
## lm(formula = mpg ~ ., data = auto)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.5903 -2.1565 -0.1169  1.8690 13.0604
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -17.218435   4.644294  -3.707  0.00024 ***
## cylinders    -0.493376   0.323282  -1.526  0.12780
## displacement  0.019896   0.007515  2.647  0.00844 **
## horsepower   -0.016951   0.013787  -1.230  0.21963
## weight       -0.006474   0.000652  -9.929 < 2e-16 ***
## acceleration  0.080576   0.098845  0.815  0.41548
## year         0.750773   0.050973  14.729 < 2e-16 ***
## origin       1.426141   0.278136  5.127  4.67e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.328 on 384 degrees of freedom
## Multiple R-squared:  0.8215, Adjusted R-squared:  0.8182
## F-statistic: 252.4 on 7 and 384 Df, p-value: < 2.2e-16
```

# Lets refit using only significant values seen in summary  
anova(fit, lm(mpg~.-cylinders-horsepower-acceleration, data = auto))

	Res.Df <dbl>	RSS <dbl>	Df <dbl>	Sum of Sq <dbl>	F <dbl>	Pr(>F) <dbl>
1	384	4252.213	NA	NA	NA	NA
2	387	4332.729	-3	-80.51617	2.423696	0.06542561
2 rows						

```
# we reject the null saying they are equal thus we use the reduced model
fit = lm(mpg~.-cylinders-horsepower-acceleration, data = auto)
lm.ar2 = summary(fit)$adj.r.squared
```

### 3(c)

We use best subset variable selection method here to and use R-square alongside  
Adjusted R-square to determine our model.

```
# best subset method

totpred <- ncol(auto) - 1
fit.full <- regsubsets(mpg~., auto, nvmax = totpred)
fit.summary <- summary(fit.full)
fit.summary
```

```
## Subset selection object
## Call: regsubsets.formula(mpg ~ ., auto, nvmax = totpred)
## 7 Variables (and intercept)
## Forced in Forced out
## cylinders FALSE FALSE
## displacement FALSE FALSE
## horsepower FALSE FALSE
## weight FALSE FALSE
## acceleration FALSE FALSE
## year FALSE FALSE
## origin FALSE FALSE
## 1 subsets of each size up to 7
## Selection Algorithm: exhaustive
## cylinders displacement horsepower weight acceleration year origin
## 1 ( 1 ) " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " " " " " " "
## 4 ( 1 ) " " " " " " " " " " " "
## 5 ( 1 ) " " " " " " " " " " " "
## 6 ( 1 ) " " " " " " " " " " " "
## 7 ( 1 ) " " " " " " " " " " " "
```

```

best.ar2 = mean(fit.summary$adjr2)

# Plot model fit measures for best model of each size against size
par(mfrow = c(2, 2))

# rsq
plot(fit.summary$rsq, xlab = "Number of Variables", ylab = "RSQ",
     type = "l")
point = as.numeric(which.max(fit.summary$rsq))
points(point, fit.summary$rsq[point],
       col = "red", cex = 2, pch = 8)

# Adjusted R^2
plot(fit.summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq",
     type = "l")
point = as.numeric(which.max(fit.summary$adjr2))
points(point, fit.summary$adjr2[point],
       col = "red", cex = 2, pch = 8)

# rsq says our best model is ALL predictors, this is bc it lessens the error
# with each additional predictor. Not a good indicator.
# AdjR2 says our very best is at 6, however the elbow can be seen around 2 or 3
# usually the elbow is our best bet with consideration to complexity


# Get coefficients of best model for a given size
coef(fit.full, 6)

```

```

## (Intercept)    cylinders displacement  horsepower      weight
## -15.563492306 -0.506685137  0.019269286 -0.023895029 -0.006218311
## year          origin
## 0.747515952  1.428241885

```



### 3(d)

Now we use forward stepwise selection

```

fit.fwd = regsubsets(mpg~., auto, nvmax = totpred, method = "forward")
fit.summary <- summary(fit.fwd)
fit.summary

```


```
## Subset selection object
## Call: regsubsets.formula(mpg ~ ., auto, nvmax = totpred, method = "forward")
## 7 Variables (and intercept)
## Forced in Forced out
## cylinders FALSE FALSE
## displacement FALSE FALSE
## horsepower FALSE FALSE
## weight FALSE FALSE
## acceleration FALSE FALSE
## year FALSE FALSE
## origin FALSE FALSE
## 1 subsets of each size up to 7
## Selection Algorithm: forward
## cylinders displacement horsepower weight acceleration year origin
## 1 (1) " " " " " " " " " " " "
## 2 (1) " " " " " " " " " " " "
## 3 (1) " " " " " " " " " " " "
## 4 (1) " " " " " " " " " " " "
## 5 (1) " " " " " " " " " " " "
## 6 (1) " " " " " " " " " " " "
## 7 (1) " " " " " " " " " " " "

fwd.ar2 = mean(fit.summary$adjr2)

# Plot model fit measures for best model of each size against size
par(mfrow = c(2, 2))
# rsq
plot(fit.summary$rsq, xlab = "Number of Variables", ylab = "RSQ",
     type = "l")
point = as.numeric(which.max(fit.summary$rsq))
points(point, fit.summary$rsq[point],
       col = "red", cex = 2, pch = 8)
# Adjusted R^2
plot(fit.summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq",
     type = "l")
point = as.numeric(which.max(fit.summary$adjr2))
points(point, fit.summary$adjr2[point],
       col = "red", cex = 2, pch = 8)

# Get coefficients of best model for a given size
coef(fit.fwd, 6)
```

	##	(Intercept)	cylinders	displacement	horsepower	weight
##	##	-15.563492306	-0.506685137	0.019269286	-0.023895029	-0.006218311
##	##					
##	##	0.747515952	1.428241885			


$$3(e)$$


Last but not least is backward stepwise selection

```
fit.bwd = regsubsets(mpg~, auto, nvmax = totpred, method = "backward")
fit.summary <- summary(fit.bwd)
fit.summary
```

```
## Subset selection object
## Call: regsubsets.formula(mpg ~ ., auto, nmax = totpred, method = "backward")
## 7 Variables (and intercept)
## Forced in Forced out
## cylinders FALSE FALSE
## displacement FALSE FALSE
## horsepower FALSE FALSE
## weight FALSE FALSE
## acceleration FALSE FALSE
## year FALSE FALSE
## origin FALSE FALSE
## 1 subsets of each size up to 7
## Selection Algorithm: backward
## cylinders displacement horsepower weight acceleration year origin
## 1 ( 1 ) " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " " " " " " " " " " "
## 4 ( 1 ) " " " " " " " " " " " " " " " "
## 5 ( 1 ) " " " " " " " " " " " " " " " "
## 6 ( 1 ) " " " " " " " " " " " " " " " "
## 7 ( 1 ) " " " " " " " " " " " " " " " "
```

```
bwd.ar2 = mean(fit.summary$adjr2)
# Plot model fit measures for best model of each size against size
par(mfrow = c(2, 2))
# rsq
plot(fit.summary$rsq, xlab = "Number of Variables", ylab = "rsq",
     type = "l")
point = as.numeric(which.max(fit.summary$rsq))
points(point, fit.summary$rsq[point],
       col = "red", cex = 2, pch = 8)
# Adjusted R^2
plot(fit.summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq",
     type = "l")
point = as.numeric(which.max(fit.summary$adjr2))
points(point, fit.summary$adjr2[point],
       col = "red", cex = 2, pch = 8)
# Get coefficients of best model for a given size
coef(fit.bwd, 6)
```

	(Intercept)	cylinders	displacement	horsepower	weight
##	-15.563492306	-0.506685137	0.019269286	-0.023895029	-0.006218311
##	year	origin			
##	0.747515952	1.428241885			



### 3(f)

Now we compare each model against a LOOCV on the data to see which model is closest to the LOOCV adjusted R-square  
Interestingly all subest methods resulted in the same model and our quick and easy MLR model is the best of the 4

```
# we compare each model using LOOCV from caret
library(caret)
set.seed(1)
tc <- trainControl(method = "loocv", verboseIter = F)
model <- train(mpg~., auto,
               method="lm", trControl = tc)
model
```

```
## Linear Regression
##
## 392 samples
## 7 predictor
##
## No pre-processing
## Resampling: Leave-One-Out Cross-Validation
## Summary of sample sizes: 391, 391, 391, 391, 391, ...
## Resampling results:
##
## RMSE      Rsquared   MAE
## 2.556369   NaN        2.556369
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

```
loocv.ar2 = summary(model)$adj.r.squared
rbind(loocv.ar2, lm.ar2, best.ar2, fwd.ar2, bwd.ar2)
```

```
##
##      [,1]
## loocv.ar2 0.8182238
## lm.ar2    0.8162176
## best.ar2  0.7979419
## fwd.ar2   0.7979419
## bwd.ar2   0.7979419
```